

Name: XXX

Position: Software Engineer, Machine Learning – Off-Search Sourcing and Relevance

Professional Summary

- Over 4 years of experience building large-scale distributed systems and online ad serving platforms
- Proficient in Python, Java, and Scala; experienced with Spark and Kafka for data pipelines
- Skilled in end-to-end machine learning workflows: data preprocessing, feature engineering, model training,
- Expertise in AWS (EC2, S3, EMR), Docker, and Kubernetes for cloud infrastructure and DevOps

Core Skills

- Programming Languages: Python, Java, Scala, SQL
- Big Data: Apache Spark, Kafka, Hadoop
- Machine Learning & Data Science: TensorFlow, PyTorch, scikit-learn, MLflow
- Distributed Systems: gRPC, RESTful APIs, microservices
- Cloud & DevOps: AWS (EC2, S3, EMR), Docker, Kubernetes, Terraform
- Ad Tech: CTR/CVR prediction, Real-Time Bidding (RTB), Information Retrieval

Experience

Senior Software Engineer | ABC Tech Co., Ltd.

July 2022 – Present, Beijing

- Designed and implemented a low-latency online ad serving system handling over 100M requests daily
- Developed CTR prediction models in TensorFlow and Spark MLlib, improving prediction accuracy by 7%
- Built real-time data pipelines (Kafka → Spark Streaming → HBase) for feature computation
- Collaborated with Product and Data Science teams on A/B tests, increasing ad conversion rates by 5%

Software Engineer | DEF Information Technology Ltd.

June 2019 – June 2022, Shanghai

- Developed an offline feature computation platform for Sponsored Products, processing hundreds of billions
- Implemented scalable ETL jobs in Python and Java with 200+ parallelism, reducing job runtime by 60%
- Contributed to multiple iterations of ad ranking and filtering algorithms to enhance ad relevance
- Maintained microservices including interface design, code reviews, and CI/CD pipelines

Projects

Off-Search Ad Relevance Engine

- Description: Engine to serve relevant ads on non-search pages (Product Detail Page, Home Page) on Ama
- Stack: TensorFlow, Spark, Kafka, AWS Lambda, Docker
- Impact: Achieved a 6% uplift in off-search CTR and an 8% revenue increase post-launch

Real-Time CTR Prediction Pipeline

- Description: Real-time feature computation and model inference pipeline using Kafka and Spark Streaming
- Stack: Python, Spark Streaming, gRPC, Kubernetes
- Impact: Maintained end-to-end latency under 50ms with 24/7 stable operation

Education

Bachelor of Science in Computer Science | Peking University

September 2015 – June 2019

- Relevant Coursework: Algorithms & Data Structures, Distributed Systems, Machine Learning, Database Sy

Certifications & Honors

- AWS Certified Solutions Architect – Associate
- TensorFlow Developer Certificate
- “Engineer of the Year” – ABC Tech Co., Ltd. (2023)

Open Source & Community

- GitHub: <https://github.com/your-username>
- Contributions: Open-sourced Spark plugins and Feature Store example implementations

Note: Replace “XXX” with a test name as needed.